

Lecture 3

CIS 341: COMPILERS

Announcements




- HW01: Hellocaml!
 - is due *tomorrow tonight* at 11:59:59pm.
- HW02: X86lite
 - Will be available soon... look for an announcement on Piazza
 - Pair-programming project
 - Simulator / Loader for x86 Assembly subset

The target architecture for CIS341

X86LITE


Intel's X86 Architecture

- 1978: Intel introduces 8086
- 1982: 80186, 80286
- 1985: 80386
- 1989: 80486 (100MHz, 1 μ m)
- 1993: Pentium
- 1995: Pentium Pro
- 1997: Pentium II/III
- 2000: Pentium 4
- 2003: Pentium M, Intel Core
- 2006: Intel Core 2
- 2008: Intel Core i3/i5/i7
- 2011: SandyBridge / IvyBridge
- 2013: Haswell
- 2014: Broadwell
- 2015: Skylake (4.2GHz, 14nm)
- AMD has a parallel line of processors



X86 Evolution & Moore's Law

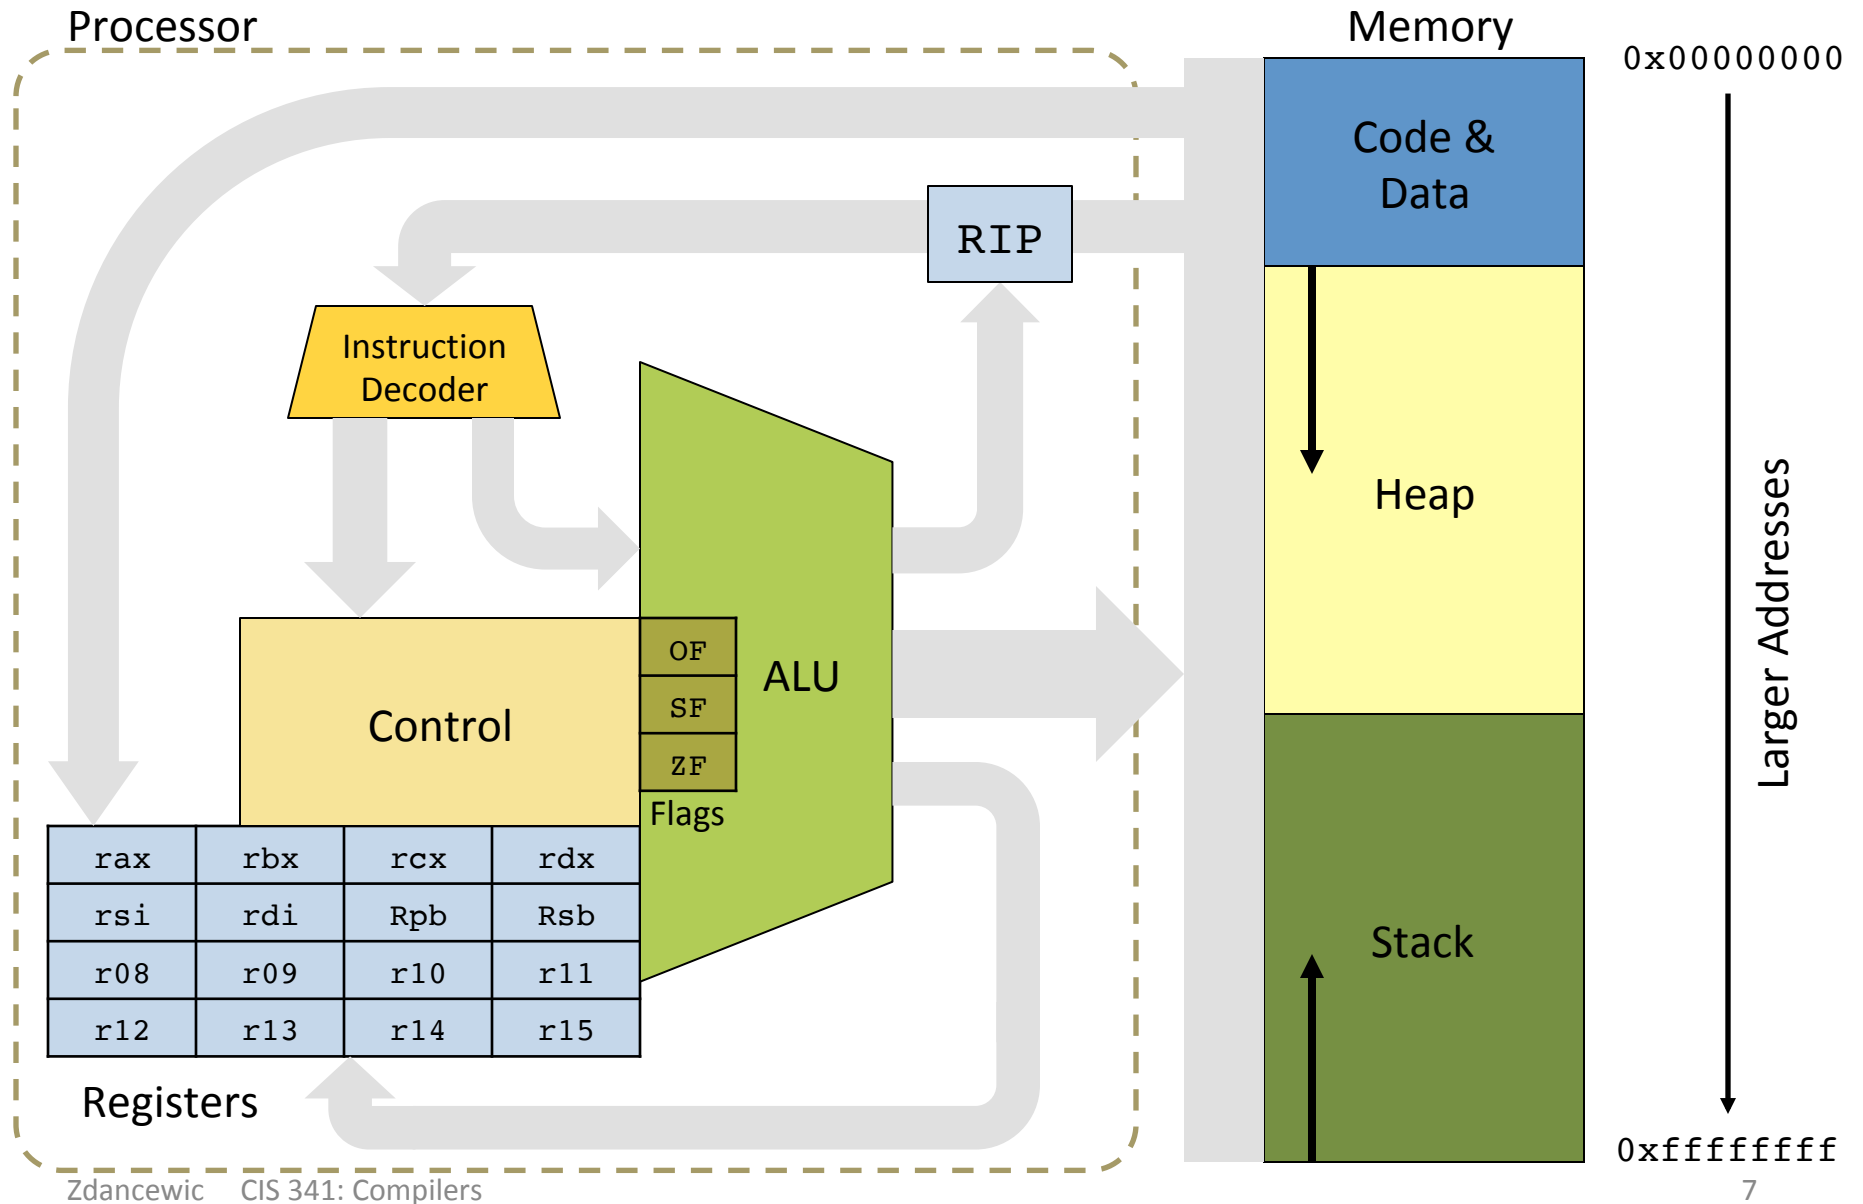
Intel Processor Transistor Count



X86 vs. X86lite

- X86 assembly is *very* complicated:
 - 8-, 16-, 32-, 64-bit values + floating points, etc.
 - Intel 64 and IA 32 architectures have a *huge* number of functions
 - “CISC” complex instructions
 - Machine code: instructions range in size from 1 byte to 17 bytes
 - Lots of hold-over design decisions for backwards compatibility
 - Hard to understand, there is a large book about optimizations at just the instruction-selection level
- X86lite is a *very* simple subset of X86:
 - Only 64 bit signed integers (no floating point, no 16bit, no ...)
 - Only about 20 instructions
 - Sufficient as a target language for general-purpose computing

X86 Schematic



X86lite Machine State: Registers

- Register File: 16 64-bit registers
 - **rax** general purpose accumulator
 - **rbx** base register, pointer to data
 - **rcx** counter register for strings & loops
 - **rdx** data register for I/O
 - **rsi** pointer register, string source register
 - **rdi** pointer register, string destination register
 - **rbp** base pointer, points to the stack frame
 - **rsp** stack pointer, points to the top of the stack
 - **R08–r15** general purpose registers
- **rip** a “virtual” register, points to the current instruction
 - **rip** is manipulated only indirectly via jumps and return.

Simplest instruction: `mov`

- `movq SRC, DEST` copy SRC into DEST
- Here, DEST and SRC are *operands*
- DEST is treated as a *location*
 - A location can be a register or a memory address
- SRC is treated as a *value*
 - A value is the *contents* of a register or memory address
 - A value can also be an *immediate* (constant) or a label
- `movq $4, %rax` // move the 64-bit immediate value 4 into `rax`
- `movq %rbx, %rax` // move the contents of `rbx` into `rax`

A Note About Instruction Syntax

- X86 presented in *two* common syntax formats

- AT&T notation: *source before* destination

- Prevalent in the Unix/Mac ecosystems
- Immediate values prefixed with '\$'
- Registers prefixed with '%'
- Mnemonic suffixes: `movq` vs. `mov`
 - `q` = quadword (4 words)
 - `l` = long (2 words)
 - `w` = word
 - `b` = byte

```
movq $5, %rax
```

```
movl $5, %eax
```

src dest

Note: X86lite uses the AT&T notation and the 64-bit only version of the instructions and registers.

- Intel notation: *destination before* source

- Used in the Intel specification / manuals
- Prevalent in the Windows ecosystem
- Instruction variant determined by register name

```
mov rax, 5
```

```
mov eax, 5
```

dest src

X86lite Arithmetic instructions

- `negq DEST` two's complement negation
- `addq SRC, DEST` $DEST \leftarrow DEST + SRC$
- `subq SRC, DEST` $DEST \leftarrow DEST - SRC$
- `imulq SRC, Reg` $Reg \leftarrow Reg * SRC$ (truncated 128-bit mult.)

Examples as written in:

```
addq %rbx, %rax    // rax ← rax + rbx
subq $4, rsp       // rsp ← rsp - 4
```

- Note: `Reg` (in `imulq`) must be a register, not a memory address

X86lite Logic/Bit manipulation Operations

- `notq DEST` logical negation
- `andq SRC, DEST` $DEST \leftarrow DEST \ \&\& \ SRC$
- `orq SRC, DEST` $DEST \leftarrow DEST \ || \ SRC$
- `xorq SRC, DEST` $DEST \leftarrow DEST \ xor \ SRC$

- `sarq Amt, DEST` $DEST \leftarrow DEST \gg amt$ (arithmetic shift right)
- `shlq Amt, DEST` $DEST \leftarrow DEST \ll amt$ (arithmetic shift left)
- `shrq Amt, DEST` $DEST \leftarrow DEST \ggg amt$ (bitwise shift right)

X86 Operands

- Operands are the values operated on by the assembly instructions
- Imm 64-bit literal signed integer “immediate”
- Lbl a “label” representing a machine address
the assembler/linker/loader resolve labels
- Reg One of the 16 registers, the value of a register is
its contents
- Ind [base:Reg][index:Reg,scale:int32][disp]
machine address (see next slide)

X86 Addressing

- In general, there are three components of an indirect address
 - **Base:** a machine address stored in a register
 - **Index * scale:** a variable offset from the base
 - **Disp:** a constant offset (displacement) from the base
- $\text{addr(ind)} = \text{Base} + [\text{Index} * \text{scale}] + \text{Disp}$
 - When used as a *location*, ind denotes the address addr(ind)
 - When used as a *value*, ind denotes $\text{Mem}[\text{addr(ind)}]$, the contents of the memory address
- Example: $-4(\%rsp)$ denotes address: $rsp - 4$
- Example: $(\%rax, \%rcx, 4)$ denotes address: $rax + 4 * rcx$
- Example: $12(\%rax, \%rcx, 4)$ denotes address: $rax + 4 * rcx + 12$
- Note: Index cannot be `rsp`

Note: X86lite does not need this full generality. It does not use index * scale.

X86lite Memory Model

- The X86lite memory consists of 2^{64} bytes numbered `0x00000000` through `0xffffffff`.
- X86lite treats the memory as consisting of 64-bit (8-byte) quadwords.
- Therefore: legal X86lite memory addresses consist of 64-bit, quadword-aligned pointers.
 - All memory addresses are evenly divisible by 8
- `leaq Ind, DEST` $\text{DEST} \leftarrow \text{addr}(\text{Ind})$ loads a pointer into DEST
- By convention, there is a stack that grows from high addresses to low addresses
- The register `rsp` points to the top of the stack
 - `pushq SRC` $\text{rsp} \leftarrow \text{rsp} - 8; \text{Mem}[\text{rsp}] \leftarrow \text{SRC}$
 - `popq DEST` $\text{DEST} \leftarrow \text{Mem}[\text{rsp}]; \text{rsp} \leftarrow \text{rsp} + 8$

X86lite State: Condition Flags & Codes

- X86 instructions set flags as a side effect
- X86lite has only 3 flags:
 - OF: “**overflow**” set when the result is too big/small to fit in 64-bit reg.
 - SF: “**sign**” set to the sign of the result (0=positive, 1 = negative)
 - ZF: “**zero**” set when the result is 0
- From these flags, we can define *Condition Codes*
 - To compare SRC1 and SRC2, compute SRC1 – SRC2 to set the flags
 - **e** equality holds when ZF is set
 - **ne** inequality holds when (not ZF)
 - **g** greater than holds when (not ZF) and (not SF)
 - **l** less than holds when SF <> OF
 - Equivalently: ((SF && not OF) || (not SF && OF))
 - **ge** greater or equal holds when (not SF)
 - **le** than or equal holds when SF <> OF or ZF

Code Blocks & Labels

- X86 assembly code is organized into *labeled blocks*:

```
label1:  
    <instruction>  
    <instruction>  
    ...  
    <instruction>  
  
label2:  
    <instruction>  
    <instruction>  
    ...  
    <instruction>
```

- Labels indicate code locations that can be jump targets (either through conditional branch instructions or function calls).
- Labels are translated away by the linker and loader – instructions live in the heap in the “code segment”
- An X86 program begins executing at a designated code label (usually “main”).

Conditional Instructions

- `cmpq SRC1, SRC2` Compute $SRC2 - SRC1$, set condition flags
- `setbCC DEST` $DEST's\ lower\ byte \leftarrow$ if CC then 1 else 0
- `jCC SRC` $rip \leftarrow$ if CC then SRC else fallthrough
- Example:

```
cmpq %rcx, %rax  
je __true1b1
```

Compare `rax` to `ecx`
If `rax = rcx` then jump to `__true1b1`

Jumps, Call and Return

- `jmp SRC` `rip` \leftarrow SRC Jump to location in SRC
- `callq SRC` Push `rip`; `rip` \leftarrow SRC
 - Call a procedure: Push the program counter to the stack (decrementing `rsp`) and then jump to the machine instruction at the address given by SRC.
- `retq` Pop into `rip`
 - Return from a procedure: Pop the current top of the stack into `rip` (incrementing `rsp`).
 - This instruction effectively jumps to the address at the top of the stack

See file: x86.ml

IMPLEMENTING X86LITE

See: runtime.c

DEMO: HANDCODING X86LITE

Compiling, Linking, Running

- To use hand-coded X86:
 1. Compile main.ml (or something like it) to either native or bytecode
 2. Run it, redirecting the output to some .s file, e.g.:
`./handcoded.native >> test.s`
 3. Use gcc to compile & link with runtime.c:
`gcc -o test runtime.c test.s`
 4. You should be able to run the resulting executable:
`./test`
- If you want to debug in gdb:
 - Call gcc with the `-g` flag too